

House Price Prediction Project

By Shaikh Abdul Quddus

Libraries

```
In [18]: import pandas as pd
from sklearn.preprocessing import LabelEncoder, StandardScaler, OneHotEncoder

from sklearn.feature_selection import SequentialFeatureSelector
from sklearn.linear_model import LinearRegression

from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer

from sklearn.model_selection import train_test_split

from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score
```

Data Gathering

```
In [4]: df=pd.read_csv('E:/Shaikh Quddus/Classes Recordings/project 1/training_set.csv')
```

EDA

```
In [5]: df.isna().sum()
```

```
Out[5]: Id                0
MSSubClass              0
MSZoning                0
LotFrontage            259
LotArea                0
...
MoSold                  0
YrSold                  0
SaleType                0
SaleCondition           0
SalePrice               0
Length: 81, dtype: int64
```

```
In [6]: df.duplicated()
```

```
Out[6]: 0      False
        1      False
        2      False
        3      False
        4      False
        ...
        1455   False
        1456   False
        1457   False
        1458   False
        1459   False
        Length: 1460, dtype: bool
```

```
In [8]: df.nunique()
```

```
Out[8]: Id                1460
        MSSubClass         15
        MSZoning           5
        LotFrontage       110
        LotArea           1073
        ...
        MoSold             12
        YrSold             5
        SaleType           9
        SaleCondition       6
        SalePrice         663
        Length: 81, dtype: int64
```

```
In [9]: df.info()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 81 columns):

#	Column	Non-Null Count	Dtype
0	Id	1460 non-null	int64
1	MSSubClass	1460 non-null	int64
2	MSZoning	1460 non-null	object
3	LotFrontage	1201 non-null	float64
4	LotArea	1460 non-null	int64
5	Street	1460 non-null	object
6	Alley	91 non-null	object
7	LotShape	1460 non-null	object
8	LandContour	1460 non-null	object
9	Utilities	1460 non-null	object
10	LotConfig	1460 non-null	object
11	LandSlope	1460 non-null	object
12	Neighborhood	1460 non-null	object
13	Condition1	1460 non-null	object
14	Condition2	1460 non-null	object
15	BldgType	1460 non-null	object
16	HouseStyle	1460 non-null	object
17	OverallQual	1460 non-null	int64
18	OverallCond	1460 non-null	int64
19	YearBuilt	1460 non-null	int64
20	YearRemodAdd	1460 non-null	int64
21	RoofStyle	1460 non-null	object
22	RoofMatl	1460 non-null	object
23	Exterior1st	1460 non-null	object
24	Exterior2nd	1460 non-null	object
25	MasVnrType	1452 non-null	object
26	MasVnrArea	1452 non-null	float64
27	ExterQual	1460 non-null	object
28	ExterCond	1460 non-null	object
29	Foundation	1460 non-null	object
30	BsmtQual	1423 non-null	object
31	BsmtCond	1423 non-null	object
32	BsmtExposure	1422 non-null	object
33	BsmtFinType1	1423 non-null	object
34	BsmtFinSF1	1460 non-null	int64
35	BsmtFinType2	1422 non-null	object
36	BsmtFinSF2	1460 non-null	int64
37	BsmtUnfSF	1460 non-null	int64
38	TotalBsmtSF	1460 non-null	int64
39	Heating	1460 non-null	object
40	HeatingQC	1460 non-null	object
41	CentralAir	1460 non-null	object
42	Electrical	1459 non-null	object
43	1stFlrSF	1460 non-null	int64
44	2ndFlrSF	1460 non-null	int64
45	LowQualFinSF	1460 non-null	int64
46	GrLivArea	1460 non-null	int64
47	BsmtFullBath	1460 non-null	int64
48	BsmtHalfBath	1460 non-null	int64
49	FullBath	1460 non-null	int64
50	HalfBath	1460 non-null	int64
51	BedroomAbvGr	1460 non-null	int64
52	KitchenAbvGr	1460 non-null	int64
53	KitchenQual	1460 non-null	object
54	TotRmsAbvGrd	1460 non-null	int64
55	Functional	1460 non-null	object
56	Fireplaces	1460 non-null	int64
57	FireplaceQu	770 non-null	object
58	GarageType	1379 non-null	object

```

59 GarageYrBlt      1379 non-null float64
60 GarageFinish     1379 non-null object
61 GarageCars       1460 non-null int64
62 GarageArea       1460 non-null int64
63 GarageQual       1379 non-null object
64 GarageCond       1379 non-null object
65 PavedDrive       1460 non-null object
66 WoodDeckSF       1460 non-null int64
67 OpenPorchSF      1460 non-null int64
68 EnclosedPorch    1460 non-null int64
69 3SsnPorch        1460 non-null int64
70 ScreenPorch      1460 non-null int64
71 PoolArea         1460 non-null int64
72 PoolQC           7 non-null object
73 Fence            281 non-null object
74 MiscFeature      54 non-null object
75 MiscVal          1460 non-null int64
76 MoSold           1460 non-null int64
77 YrSold           1460 non-null int64
78 SaleType         1460 non-null object
79 SaleCondition    1460 non-null object
80 SalePrice        1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

Define x and y

```

In [13]: x=df.drop(['SalePrice','Id'],axis=1)
         y=df['SalePrice']

```

Preprocessing for Feature Selection

Missing values handling

```

In [14]: for i in x.columns:
         if x[i].dtypes=='object':
             x[i]=x[i].fillna(x[i].mode()[0])
         else:
             x[i]=x[i].fillna(x[i].median())

```

```

In [16]: x.info()

```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1460 entries, 0 to 1459

Data columns (total 79 columns):

#	Column	Non-Null Count	Dtype
0	MSSubClass	1460 non-null	int64
1	MSZoning	1460 non-null	object
2	LotFrontage	1460 non-null	float64
3	LotArea	1460 non-null	int64
4	Street	1460 non-null	object
5	Alley	1460 non-null	object
6	LotShape	1460 non-null	object
7	LandContour	1460 non-null	object
8	Utilities	1460 non-null	object
9	LotConfig	1460 non-null	object
10	LandSlope	1460 non-null	object
11	Neighborhood	1460 non-null	object
12	Condition1	1460 non-null	object
13	Condition2	1460 non-null	object
14	BldgType	1460 non-null	object
15	HouseStyle	1460 non-null	object
16	OverallQual	1460 non-null	int64
17	OverallCond	1460 non-null	int64
18	YearBuilt	1460 non-null	int64
19	YearRemodAdd	1460 non-null	int64
20	RoofStyle	1460 non-null	object
21	RoofMatl	1460 non-null	object
22	Exterior1st	1460 non-null	object
23	Exterior2nd	1460 non-null	object
24	MasVnrType	1460 non-null	object
25	MasVnrArea	1460 non-null	float64
26	ExterQual	1460 non-null	object
27	ExterCond	1460 non-null	object
28	Foundation	1460 non-null	object
29	BsmtQual	1460 non-null	object
30	BsmtCond	1460 non-null	object
31	BsmtExposure	1460 non-null	object
32	BsmtFinType1	1460 non-null	object
33	BsmtFinSF1	1460 non-null	int64
34	BsmtFinType2	1460 non-null	object
35	BsmtFinSF2	1460 non-null	int64
36	BsmtUnfSF	1460 non-null	int64
37	TotalBsmtSF	1460 non-null	int64
38	Heating	1460 non-null	object
39	HeatingQC	1460 non-null	object
40	CentralAir	1460 non-null	object
41	Electrical	1460 non-null	object
42	1stFlrSF	1460 non-null	int64
43	2ndFlrSF	1460 non-null	int64
44	LowQualFinSF	1460 non-null	int64
45	GrLivArea	1460 non-null	int64
46	BsmtFullBath	1460 non-null	int64
47	BsmtHalfBath	1460 non-null	int64
48	FullBath	1460 non-null	int64
49	HalfBath	1460 non-null	int64
50	BedroomAbvGr	1460 non-null	int64
51	KitchenAbvGr	1460 non-null	int64
52	KitchenQual	1460 non-null	object
53	TotRmsAbvGrd	1460 non-null	int64
54	Functional	1460 non-null	object
55	Fireplaces	1460 non-null	int64
56	FireplaceQu	1460 non-null	object
57	GarageType	1460 non-null	object
58	GarageYrBlt	1460 non-null	float64

```

59 GarageFinish    1460 non-null object
60 GarageCars      1460 non-null int64
61 GarageArea      1460 non-null int64
62 GarageQual      1460 non-null object
63 GarageCond      1460 non-null object
64 PavedDrive      1460 non-null object
65 WoodDeckSF      1460 non-null int64
66 OpenPorchSF     1460 non-null int64
67 EnclosedPorch   1460 non-null int64
68 3SsnPorch       1460 non-null int64
69 ScreenPorch     1460 non-null int64
70 PoolArea        1460 non-null int64
71 PoolQC          1460 non-null object
72 Fence           1460 non-null object
73 MiscFeature     1460 non-null object
74 MiscVal         1460 non-null int64
75 MoSold          1460 non-null int64
76 YrSold          1460 non-null int64
77 SaleType        1460 non-null object
78 SaleCondition   1460 non-null object
dtypes: float64(3), int64(33), object(43)
memory usage: 901.2+ KB

```

```
In [17]: cat=x.select_dtypes(include='object')
con=x.select_dtypes(exclude='object')
```

```
In [19]: le=LabelEncoder()
```

```
In [20]: cat1=cat.apply(le.fit_transform)
cat1
```

```
Out[20]:
```

	MSZoning	Street	Alley	LotShape	LandContour	Utilities	LotConfig	LandSlope	Neighbor
0	3	1	0	3	3	0	4	0	
1	3	1	0	3	3	0	2	0	
2	3	1	0	0	3	0	4	0	
3	3	1	0	0	3	0	0	0	
4	3	1	0	0	3	0	2	0	
...	
1455	3	1	0	3	3	0	4	0	
1456	3	1	0	3	3	0	4	0	
1457	3	1	0	3	3	0	4	0	
1458	3	1	0	3	3	0	4	0	
1459	3	1	0	3	3	0	4	0	

1460 rows × 43 columns

```
In [21]: ss=StandardScaler()
```

```
In [22]: con1=pd.DataFrame(ss.fit_transform(con),columns=ss.get_feature_names_out())
```

```
In [23]: con1
```

Out[23]:

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	I
0	0.073375	-0.220875	-0.207142	0.651479	-0.517200	1.050994	0.878668	
1	-0.872563	0.460320	-0.091886	-0.071836	2.179628	0.156734	-0.429577	
2	0.073375	-0.084636	0.073480	0.651479	-0.517200	0.984752	0.830215	
3	0.309859	-0.447940	-0.096897	0.651479	-0.517200	-1.863632	-0.720298	
4	0.073375	0.641972	0.375148	1.374795	-0.517200	0.951632	0.733308	
...	
1455	0.073375	-0.357114	-0.260560	-0.071836	-0.517200	0.918511	0.733308	
1456	-0.872563	0.687385	0.266407	-0.071836	0.381743	0.222975	0.151865	
1457	0.309859	-0.175462	-0.147810	0.651479	3.078570	-1.002492	1.024029	
1458	-0.872563	-0.084636	-0.080160	-0.795151	0.381743	-0.704406	0.539493	
1459	-0.872563	0.233255	-0.058112	-0.795151	0.381743	-0.207594	-0.962566	

1460 rows × 36 columns

In [24]: `x1=con1.join(cat1)`
`x1`

Out[24]:

	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt	YearRemodAdd	I
0	0.073375	-0.220875	-0.207142	0.651479	-0.517200	1.050994	0.878668	
1	-0.872563	0.460320	-0.091886	-0.071836	2.179628	0.156734	-0.429577	
2	0.073375	-0.084636	0.073480	0.651479	-0.517200	0.984752	0.830215	
3	0.309859	-0.447940	-0.096897	0.651479	-0.517200	-1.863632	-0.720298	
4	0.073375	0.641972	0.375148	1.374795	-0.517200	0.951632	0.733308	
...	
1455	0.073375	-0.357114	-0.260560	-0.071836	-0.517200	0.918511	0.733308	
1456	-0.872563	0.687385	0.266407	-0.071836	0.381743	0.222975	0.151865	
1457	0.309859	-0.175462	-0.147810	0.651479	3.078570	-1.002492	1.024029	
1458	-0.872563	-0.084636	-0.080160	-0.795151	0.381743	-0.704406	0.539493	
1459	-0.872563	0.233255	-0.058112	-0.795151	0.381743	-0.207594	-0.962566	

1460 rows × 79 columns

Feature Selection

In [25]: `lr=LinearRegression()`

In [156... `sfs=SequentialFeatureSelector(lr,n_features_to_select='auto',tol=None)`

```
In [157... sfs.fit(x1,y)
```

```
Out[157]: SequentialFeatureSelector
          estimator: LinearRegression
              LinearRegression
```

```
In [158... cols=sfs.get_feature_names_out()
```

```
In [159... cols
```

```
Out[159]: array(['MSSubClass', 'LotArea', 'OverallQual', 'OverallCond', 'YearBuilt',
        'MasVnrArea', 'BsmtFinSF1', 'GrLivArea', 'BsmtFullBath',
        'KitchenAbvGr', 'TotRmsAbvGrd', 'Fireplaces', 'GarageCars',
        'WoodDeckSF', 'OpenPorchSF', 'ScreenPorch', 'PoolArea', 'YrSold',
        'Street', 'LandContour', 'Utilities', 'Neighborhood', 'BldgType',
        'HouseStyle', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'MasVnrType',
        'ExterQual', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'HeatingQC',
        'KitchenQual', 'Functional', 'GarageCond', 'PavedDrive', 'Fence',
        'MiscFeature'], dtype=object)
```

Final Dataset

```
In [160... x2=pd.DataFrame(df,columns=cols)
```

```
In [161... x2
```

```
Out[161]:
```

	MSSubClass	LotArea	OverallQual	OverallCond	YearBuilt	MasVnrArea	BsmtFinSF1	GrLivA
0	60	8450	7	5	2003	196.0	706	1
1	20	9600	6	8	1976	0.0	978	1
2	60	11250	7	5	2001	162.0	486	1
3	70	9550	7	5	1915	0.0	216	1
4	60	14260	8	5	2000	350.0	655	2
...
1455	60	7917	6	5	1999	0.0	0	1
1456	20	13175	6	6	1978	119.0	790	2
1457	70	9042	7	9	1941	0.0	275	2
1458	20	9717	5	6	1950	0.0	49	1
1459	20	9937	5	6	1965	0.0	830	1

1460 rows × 39 columns

Preprocessing


```
In [162... cat=[]
con=[]

for i in x2.columns:
    if x2[i].dtypes=='object':
        cat.append(i)
    else:
        con.append(i)
```

```
In [163... cat
```

```
Out[163]: ['Street',
'LandContour',
'Utilities',
'Neighborhood',
'BldgType',
'HouseStyle',
'RoofStyle',
'RoofMatl',
'Exterior1st',
'MasVnrType',
'ExterQual',
'BsmtQual',
'BsmtCond',
'BsmtExposure',
'HeatingQC',
'KitchenQual',
'Functional',
'GarageCond',
'PavedDrive',
'Fence',
'MiscFeature']
```

```
In [164... con
```

```
Out[164]: ['MSSubClass',
'LotArea',
'OverallQual',
'OverallCond',
'YearBuilt',
'MasVnrArea',
'BsmtFinSF1',
'GrLivArea',
'BsmtFullBath',
'KitchenAbvGr',
'TotRmsAbvGrd',
'Fireplaces',
'GarageCars',
'WoodDeckSF',
'OpenPorchSF',
'ScreenPorch',
'PoolArea',
'YrSold']
```

```
In [165... num_pipe=Pipeline(steps=[('impute',SimpleImputer(strategy='median')),('scaler',Star
cat_pipe=Pipeline(steps=[('Impute',SimpleImputer(strategy='most_frequent')),('encoc
```

```
In [166... pre=ColumnTransformer([('num_pipe',num_pipe,con),('cat_pipe',cat_pipe,cat)])
```

```
In [167... x3=pd.DataFrame(pre.fit_transform(x2).toarray(),columns=pre.get_feature_names_out()
x3
```

Out[167]:

	num_pipe_MSSubClass	num_pipe_LotArea	num_pipe_OverallQual	num_pipe_OverallCond
0	0.073375	-0.207142	0.651479	-0.517200
1	-0.872563	-0.091886	-0.071836	2.179628
2	0.073375	0.073480	0.651479	-0.517200
3	0.309859	-0.096897	0.651479	-0.517200
4	0.073375	0.375148	1.374795	-0.517200
...
1455	0.073375	-0.260560	-0.071836	-0.517200
1456	-0.872563	0.266407	-0.071836	0.381743
1457	0.309859	-0.147810	0.651479	3.078570
1458	-0.872563	-0.080160	-0.795151	0.381743
1459	-0.872563	-0.058112	-0.795151	0.381743

1460 rows × 145 columns

Splitting data into Train and Test

In [176... `x_train,x_test,y_train,y_test=train_test_split(x3,y,test_size=0.2,random_state=18)`

Model

In [177... `lr.fit(x_train,y_train)`

Out[177]:

LinearRegression

LinearRegression()

Training Data Evaluation

In [178... `y_pred_train=lr.predict(x_train)`

```

mse=mean_squared_error(y_pred_train,y_train)
print('MSE:',mse)

rmse=mse**0.5
print('RMSE:',rmse)

mae=mean_absolute_error(y_pred_train,y_train)
print('MAE:',mae)

R=round(r2_score(y_pred_train,y_train),2)
print('R2:',R)

```

MSE: 593642586.3535959
RMSE: 24364.78168081126
MAE: 15167.007705479453
R2: 0.9

Testing Data Evaluation

```
In [179... y_pred=lr.predict(x_test)

mse1=mean_squared_error(y_pred,y_test)
print('Mse:',mse1)

rmse1=mse1**0.5
print('RMSE:',rmse1)

mae1=mean_absolute_error(y_pred,y_test)
print('MAE:',mae1)

R1=r2_score(y_pred,y_test)
print('R2:',R1)
```

Mse: 7.396956465369425e+28
RMSE: 271973463142443.8
MAE: 19464887836652.95
R2: -0.0017301716410542678

Regularization

```
In [103... from sklearn.linear_model import Ridge
```

```
In [180... ra=Ridge()
```

```
In [181... ra.fit(x_train, y_train)
```

```
Out[181]: ▼ Ridge
Ridge()
```

```
In [182... y_pred_train=ra.predict(x_train)

mse=mean_squared_error(y_pred_train,y_train)
print('MSE:',mse)

rmse=mse**0.5
print('RMSE:',rmse)

mae=mean_absolute_error(y_pred_train,y_train)
print('MAE:',mae)

R=r2_score(y_pred_train,y_train)
print('R2:',R)
```

MSE: 659993445.507732
RMSE: 25690.337590380786
MAE: 16031.827815828041
R2: 0.8796527814919366

```
In [183... y_pred=ra.predict(x_test)

mse1=mean_squared_error(y_pred,y_test)
print('Mse:',mse1)

rmse1=mse1**0.5
print('RMSE:',rmse1)

mae1=mean_absolute_error(y_pred,y_test)
print('MAE:',mae1)

R1=r2_score(y_pred,y_test)
print('R2:',R1)
```

```
Mse: 631833640.5118048
RMSE: 25136.30124962312
MAE: 16860.163469639647
R2: 0.8859275614048505
```

Hyperparameter Tuning

```
In [184... import numpy as np
```

```
In [185... grid={
    'alpha':np.arange(1,100,0.1)
}
```

```
In [186... from sklearn.model_selection import GridSearchCV,RandomizedSearchCV
from warnings import filterwarnings
filterwarnings('ignore')
```

```
In [187... rs=RandomizedSearchCV(ra,param_distributions=grid,cv=6)
rs.fit(x_train,y_train)
```

```
Out[187]: ▸ RandomizedSearchCV
          ▸ estimator: Ridge
              ▸ Ridge
```

```
In [188... rs.best_params_
```

```
Out[188]: {'alpha': 5.7000000000000005}
```

```
In [189... ra=Ridge(5.7000000000000005)
```

```
In [190... ra.fit(x_train, y_train)
```

```
Out[190]: ▾ Ridge
          Ridge(alpha=5.7000000000000005)
```

```
In [191... y_pred_train=ra.predict(x_train)

mse=mean_squared_error(y_pred_train,y_train)
print('MSE:',mse)
```

```
rmse=mse**0.5
print('RMSE:',rmse)

mae=mean_absolute_error(y_pred_train,y_train)
print('MAE:',mae)

R=round(r2_score(y_pred_train,y_train),2)
print('R2:',R)
```

```
MSE: 751675950.4725881
RMSE: 27416.709329760713
MAE: 16537.012170830672
R2: 0.86
```

```
In [192... y_pred=ra.predict(x_test)

mse1=mean_squared_error(y_pred,y_test)
print('Mse:',mse1)

rmse1=mse1**0.5
print('RMSE:',rmse1)

mae1=mean_absolute_error(y_pred,y_test)
print('MAE:',mae1)

R1=r2_score(y_pred,y_test)
print('R2:',R1)
```

```
Mse: 672023890.2534853
RMSE: 25923.42358280413
MAE: 16701.223788877203
R2: 0.8725162671547294
```

unseen data

```
In [194... df2=pd.read_csv('E:/Shaikh Quddus/Classes Recordings/project 1/testing_set.csv')
```

```
In [195... x_samp=pre.transform(df2).toarray()
```

```
In [196... x_samp=pd.DataFrame(x_samp,columns=pre.get_feature_names_out())
```

```
In [197... x_samp
```

Out[197]:	num_pipe_MSSubClass	num_pipe_LotArea	num_pipe_OverallQual	num_pipe_OverallCond
0	-0.872563	0.110763	-0.795151	0.381743
1	-0.872563	0.375850	-0.071836	0.381743
2	0.073375	0.332053	-0.795151	-0.517200
3	0.073375	-0.054002	-0.071836	0.381743
4	1.492282	-0.552407	1.374795	-0.517200
...
1454	2.438219	-0.859988	-1.518467	1.280685
1455	2.438219	-0.864197	-1.518467	-0.517200
1456	-0.872563	0.950423	-0.795151	1.280685
1457	0.664586	-0.007600	-0.795151	-0.517200
1458	0.073375	-0.089180	0.651479	-0.517200

1459 rows × 145 columns



```
In [198... y_pred1=ra.predict(x_samp)
y_pred1
```

```
Out[198]: array([121338.04299669, 151894.75801705, 164133.46680722, ...,
178881.60017101, 107994.00430503, 216021.19734466])
```

```
In [199... df3=df2[['Id']]]
```

```
In [200... df3['SalePrice']=y_pred1
```

```
In [201... df3
```

Out[201]:	Id	SalePrice
0	1461	121338.042997
1	1462	151894.758017
2	1463	164133.466807
3	1464	184005.576569
4	1465	211708.459400
...
1454	2915	74927.375617
1455	2916	81225.776649
1456	2917	178881.600171
1457	2918	107994.004305
1458	2919	216021.197345

1459 rows × 2 columns

Saving the predicted salesprice in a csv

```
In [203... df3.to_csv('E:/house_salesprice.csv',index=False)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```

```
In [ ]:
```